

CPSC 501 A4

Teodor Tebeica

30046038

MNISTModel.py

```
--Fit model--
Epoch 1/10
469/469 - 1s - loss: 0.2829 - accuracy: 0.9193
Epoch 2/10
469/469 - 1s - loss: 0.1220 - accuracy: 0.9645
Epoch 3/10
469/469 - 1s - loss: 0.0834 - accuracy: 0.9748
Epoch 4/10
469/469 - 1s - loss: 0.0631 - accuracy: 0.9811
Epoch 5/10
469/469 - 1s - loss: 0.0501 - accuracy: 0.9845
Epoch 6/10
469/469 - 1s - loss: 0.0410 - accuracy: 0.9877
Epoch 7/10
469/469 - 1s - loss: 0.0331 - accuracy: 0.9900
Epoch 8/10
469/469 - 1s - loss: 0.0289 - accuracy: 0.9910
Epoch 9/10
469/469 - 1s - loss: 0.0235 - accuracy: 0.9928
Epoch 10/10
469/469 - 1s - loss: 0.0211 - accuracy: 0.9935
--Evaluate model--
313/313 - 0s - loss: 0.0617 - accuracy: 0.9818
Model Loss:    0.06
Model Accuracy: 98.2%

Process finished with exit code 0
```

Results: 99.35% on training data, 98.18% on test data

To run: python MNISTModel.py

For this part of the assignment, I have had to change the layers of the neural network model as well as the optimizer, batch_size and number of epochs, as follows:

1. Keras model (now 4 layers)
 - a. Flatten
 - i. Stayed the same as the starter code
 - b. Dense

- i. `tf.keras.layers.Dense(512, activation='relu')`
 - ii. a densely connected neural network layer, with 512 outputs with ReLU activation which helps rejecting very wrong guesses giving this layer better performance
 - c. Dropout
 - i. `tf.keras.layers.Dropout(0.2)`
 - ii. speeds up the training process, forces nodes within a layer to take more or less responsibility given a probability (0.2) in this case
 - d. Dense
 - i. `tf.keras.layers.Dense(10, activation='softmax')`
 - ii. a densely connected neural network layer, with 10 outputs with softmax activation which helps being more favorable for output when more than 2 categories need to be considered.
- 2. Optimizer
 - a. Changed the optimizer to adam – seems to be more commonly used and after testing other optimizers as well, I concluded that adam delivers the most optimized performance for this model
- 3. Batch size has been increased to 128 speeds up the time spent in each epoch without losing much accuracy.
- 4. Number of epochs has been increased to 10 allowing model to train on more training data before going for the testing data.

notMNISTModel.py

Originally, I have created this model similar to the MNISTModel.py one.

The model was as follows:

- tf.keras.layers.Flatten(input_shape=(28, 28)),
- tf.keras.layers.Dense(256, activation='relu'),
- tf.keras.layers.Dropout(0.1),
- tf.keras.layers.Dense(10, activation='softmax')
- Epochs set to 10


Model works well with capital letters from the predict_test.py application.

```
--Evaluate model--  
313/313 - 0s - loss: 0.2183 - accuracy: 0.9387  
Model Loss: 0.22  
Model Accuracy: 93.9%
```

Modifications for predict files:

- Line to load your model
 - o model = tf.keras.models.load_model(sys.argv[2])
(provided you are importing tensorflow as tf)
- Line to get array of percent confidence
 - o prediction = model.predict(img)[0]
- Line to decide the index of the highest prediction
 - o predicted_label = prediction.argmax(axis=-1)

However, the model identified **incorrectly** the following image of a letter of “F” as an “I” with 51% accuracy.




After change, the model correctly identified the letter as F, with 95% accuracy.

The model used was as follows:

```
--Evaluate model--  
313/313 - 0s - loss: 0.2973 - accuracy: 0.9393  
Model Loss: 0.30  
Model Accuracy: 93.9%
```

- tf.keras.layers.Flatten(input_shape=(28, 28))
- tf.keras.layers.Dense(512, activation='relu')
- tf.keras.layers.Dropout(0.1)
- tf.keras.layers.Dense(10, activation='softmax')
- changed epochs **to equal 30**
- batch size stays the same as the original



To run: `python notMNISTModel.py`

CHDModel.py

BEFORE:

I implemented a function to split my heart.csv data randomly 90% into heart_train.csv and 10% into heart_test.csv.

I, then, followed [this](#) tutorial in order to load the csv components into the program so that they can be used with tensorflow (apply data normalization).

The original model was showing signs of **overfit**. The train datasets were evaluated at 97.8% whereas the test datasets evaluated at 65% accuracy.

The original model was as follows:

- `tf.keras.layers.DenseFeatures(categoricalColumns + numericColumns)`
- `tf.keras.layers.Dense(256, activation='relu')`
- `tf.keras.layers.Dense(256, activation='relu')`
- `tf.keras.layers.Dense(1, activation='sigmoid')`

The compilation of the model was done as follows:

- `model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])`
- 20 epochs at a size of 128 each to fit model

Results:

```
Epoch 20/20
128/128 [=====] - 0s 623us/step - loss: 0.1070 - accuracy: 0.9781
--Evaluate model--
WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor, but we
Consider rewriting this model with the Functional API.
128/128 [=====] - 0s 475us/step - loss: 1.4412 - accuracy: 0.6500
Model Loss:    1.44
Model Accuracy: 65.0%
```

AFTER:

After some tweaking of the model, I have managed to decrease its loss rate by about 3 times, reaching an accuracy of 78.2% on test data and 72.2% on train data.

I followed [this](#) tutorial.

The modified model is as follows:

- `tf.keras.layers.DenseFeatures(categoricalColumns + numericColumns)`
- `m.add(tf.keras.layers.Dense(32, kernel_regularizer=tf.keras.regularizers.l2(0.01)))`
- `tf.keras.layers.Dropout(0.5) tf.keras.layers.Dense(32, kernel_regularizer=tf.keras.regularizers.l2(0.01), activation='relu')`
- `tf.keras.layers.Dropout(0.5)`
- `tf.keras.layers.Dense(1, activation='sigmoid')`

The compilation of the model is done as follows:

- `model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])`
- 20 epochs at a size of **512 each** to fit model

Results:

```
Epoch 20/20
512/512 [=====] - 0s 522us/step - loss: 0.5471 - accuracy: 0.7297
--Evaluate model--
WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor, but we
Consider rewriting this model with the Functional API.
WARNING:tensorflow:Callbacks method `on_test_batch_end` is slow compared to the batch time (b
512/512 [=====] - 0s 446us/step - loss: 0.5116 - accuracy: 0.7395
Model Loss:    0.51
Model Accuracy: 73.9%
```

```
Epoch 20/20
512/512 [=====] - 0s 559us/step - loss: 0.5630 - accuracy: 0.7219
--Evaluate model--
WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor, but we
Consider rewriting this model with the Functional API.
512/512 [=====] - 0s 464us/step - loss: 0.4455 - accuracy: 0.7824
Model Loss:    0.45
Model Accuracy: 78.2%
```

These changes made the training data accuracy to drop for their original 97.8% to 72.2%. This outcome did positively impact the test model accuracy, raising it from 65% to 78.2%.

To run: python CHDModel.py